


MPI Processes

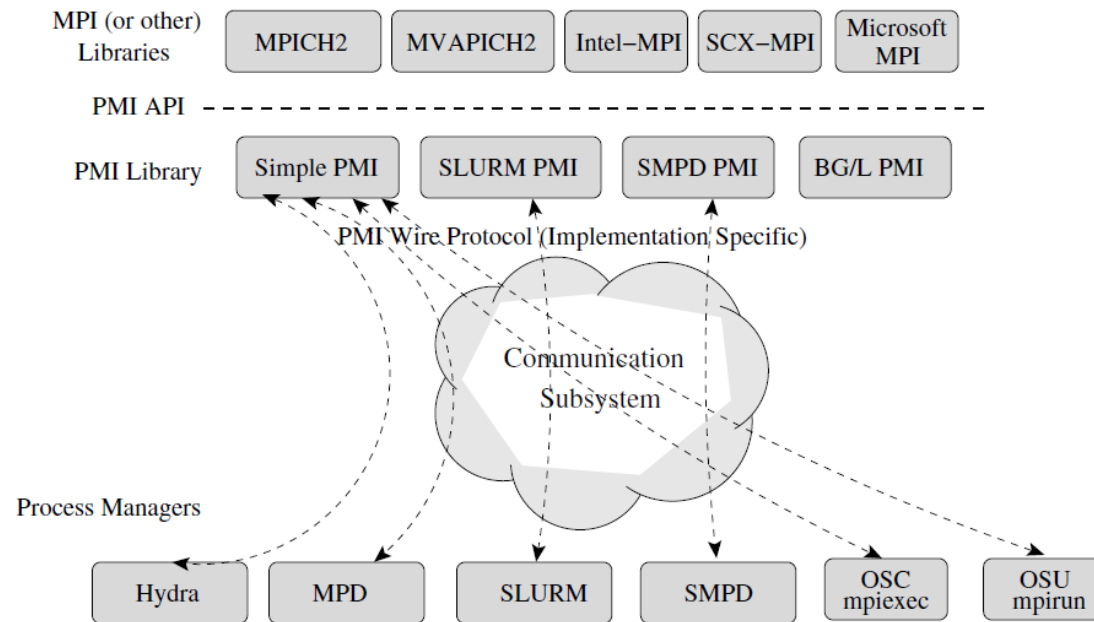
Jan 8, 2019

Hardware Model



NO centralized server/master

Process Management Setup



**Parallel program
library (e.g. MPI)**

**Process
management
interface (PMI)**

**Resource manager/
Job scheduler**

Reference

[PMI: A Scalable Parallel Process-Management Interface for Extreme-Scale Systems](#)

Internals


Process Manager

- Start and stop processes in a **scalable** way
- Setup communication channels for parallel processes
- Provide system-specific information to processes


Process Management Interface

- Processes can exchange information about peers by querying PMI
- Provides a logically centralized service for all processes in an MPI job
- Uses key-value store for process-related data

Process Launch



Process Launch




Hydra Process Manager

- A process management system for starting parallel jobs
- Uses existing daemons (viz. ssh) to start MPI processes
- Automatically detects resource managers and interacts with them
- `$ mpiexec ./app`
 - Hydra gets information about allocated resources and launches processes
- Passes environment variables from the shell on which mpiexec is launched to the launched processes

There are others – mpd, gforker, slurm, etc.

mpiexec



```
mpiexec -n 4 -hosts host1,host2,host3,host4 ./exe
```


Launch Node

`mpixexec -np 8 -hosts host1:3,host2:3,host3:3 ./exe`


```
pmalakkar 17952 17943 0 09:41 ? 00:00:00 /usr/lib/openssh/sftp-server
pmalakkar 20853 16203 0 10:20 pts/1 00:00:00 mpixexec -np 8 -hosts 172.27.19.2 3 172.27.19.3 3 172.27.19.4 3 ./IMB-MPI1 AllReduce
pmalakkar 20854 20853 0 10:20 ? 00:00:00 /users/faculty/pmalakkar/mpich-3.2.1-install/bin/hydra_pmi_proxy --control-port 172.27.19.2:46385 --rmk user --launcher ssh --demux poll --pgid 0 --retries 10 --usize -2 --proxy-id 0
pmalakkar 20855 20853 0 10:20 ? 00:00:00 /usr/bin/ssh -x 172.27.19.3 "/users/faculty/pmalakkar/mpich-3.2.1-install/bin/hydra_pmi_proxy" --control-port 172.27.19.2:46385 --rmk user --launcher ssh --demux poll --pgid 0 --retries 10 --usize -2 --proxy-id 1
pmalakkar 20856 20853 0 10:20 ? 00:00:00 /usr/bin/ssh -x 172.27.19.4 "/users/faculty/pmalakkar/mpich-3.2.1-install/bin/hydra_pmi_proxy" --control-port 172.27.19.2:46385 --rmk user --launcher ssh --demux poll --pgid 0 --retries 10 --usize -2 --proxy-id 2
pmalakkar 20857 20854 76 10:20 ? 00:00:03 ./IMB-MPI1 AllReduce
pmalakkar 20858 20854 76 10:20 ? 00:00:03 ./IMB-MPI1 AllReduce
pmalakkar 20859 20854 76 10:20 ? 00:00:03 ./IMB-MPI1 AllReduce
pmalakkar 20861 17877 0 10:20 pts/4 00:00:00 ps -aef
```

Compute Node Processes

```
pmalakar 8756 8728 0 10:18 pts/0 00:00:00 -bash
pmalakar 8759 8755 0 10:18 ? 00:00:00 /usr/lib/openssh/sftp-server
root 8781 1123 0 10:20 ? 00:00:00 sshd: pmalakar [priv]
pmalakar 8845 8781 0 10:20 ? 00:00:00 sshd: pmalakar@notty
pmalakar 8846 8845 0 10:20 ? 00:00:00 /users/faculty/pmalakar/mpich-3.2.1-install/bin/hydra_pmi_proxy
y --control-port 172.27.19.2:46385 --rmk user --launcher ssh --demux poll --pgid 0 --retries 10 --usize -2 --p
roxy-id 1
pmalakar 8847 8846 99 10:20 ? 00:00:12 ./IMB-MPI1 AllReduce
pmalakar 8848 8846 99 10:20 ? 00:00:12 ./IMB-MPI1 AllReduce
pmalakar 8849 8846 99 10:20 ? 00:00:12 ./IMB-MPI1 AllReduce
```

```
pmalakar 8838 8774 0 10:20 pts/1 00:00:00 -bash
pmalakar 8841 8837 0 10:20 ? 00:00:00 /usr/lib/openssh/sftp-server
root 8851 1250 0 10:20 ? 00:00:00 sshd: pmalakar [priv]
pmalakar 8915 8851 0 10:20 ? 00:00:00 sshd: pmalakar@notty
pmalakar 8916 8915 0 10:20 ? 00:00:00 /users/faculty/pmalakar/mpich-3.2.1-install/bin/hydra_p
mi_proxy --control-port 172.27.19.2:46385 --rmk user --launcher ssh --demux poll --pgid 0 --retries 10
--usize -2 --proxy-id 2
pmalakar 8917 8916 99 10:20 ? 00:00:14 ./IMB-MPI1 AllReduce
pmalakar 8918 8916 99 10:20 ? 00:00:14 ./IMB-MPI1 AllReduce
```

Hardware Model



What are different ways in which you can share data among distributed processes?


Message Passing Paradigm

- Message sends and receives
- Explicit communication


Communication types

- Blocking
- Non-blocking

Hardware Model



Communication Channels




- Sockets for network I/O (wire protocol in PMI)
- PMI is responsible for creating/initializing/cleanup
- MPI handles communications, progress etc.

Reading: Design and Evaluation of Nemesis, a Scalable, Low-Latency, Message-Passing Communication Subsystem by Buntinas et al.

Interconnects


Linear Array



Attributes / Parameters


- Topology
- Diameter $p-1$
- Bisection width 1
- Cost $p-1$

Ring Interconnect




- Diameter $p/2$
- Bisection width 2
- Cost p

Star Interconnect




- Diameter 2
- Bisection width 1
- Cost $p-1$

Mesh Interconnect



- Diameter $2(\sqrt{p} - 1)$
- Bisection width \sqrt{p}
- Cost $2(p - \sqrt{p})$


Torus Interconnect



What is the advantage of torus over mesh?


- Diameter $2(\sqrt{p}/2)$
- Bisection width \sqrt{p}
- Cost $2p$

Hypercube Interconnect




- Diameter $\log p$
- Bisection width $p/2$
- Cost $(p \log p)/2$

Switched Network



Crossbar Switch



Connectivity ?
Latency ?

Communication Cost

- Startup time (t_s)
- Latency (t_h)
- Bandwidth (t_w)


$$\text{Transfer time} = t_s + t_h + n/t_w$$

$$\text{Communication time} = \text{Transfer time} + \text{Overhead}$$

MPI Programming

Parallel Program Execution

- Launch MPI processes on cluster nodes
- Communication setup



MPI

- Standard for message passing
- Explicit communications
- High programming complexity
- Requires communication scope

Getting Started

Initializes
and queries
PMI

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {

    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    // Get the rank of the process
    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    // Print off a hello world message
    printf("Hello I am rank %d out of %d processes\n", rank, size);

    // Finalize the MPI environment.
    MPI_Finalize();
}
```

MPI_Init

- gather information about the parallel job
- set up internal library state
- prepare for communication

Communication Scope

Process

- belongs to a group
- identified by a rank within a group

Message

- context
- tag

A communication handle **Communicator** defines the scope

Getting Started

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {

    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    // Get the rank of the process
    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    // Print a hello world message
    printf("Hello world! I am rank %d out of %d processes\n", rank, size);

    // Clean up the MPI environment.
    MPI_Finalize();
}
```

Total
number of
processes


Rank of a
process

MPI Communication Types

Point-to-point




Collective



Basic MPI Communication

- MPI_Send




SENDER

Blocking send and receive

```
int MPI_Send (const void *buf, int count,  
MPI_Datatype datatype, int dest, int tag,  
MPI_Comm comm)
```

Tags should match

- MPI_Recv



RECEIVER

```
int MPI_Recv (void *buf, int count,  
MPI_Datatype datatype, int source, int tag,  
MPI_Comm comm, MPI_Status *status)
```

Teaching Assistants

- Dixit Kumar (dixit)
 - Kawal Preet (kawal)
 - Nirjhar Roy (nirjhar)
 - Soumya Banerjee (soumyab)
-
- Add them to your Git repo