


MPI-II

Jan 15, 2019

Recap of Basic Concepts and Terminology

- Message passing
- Node
- Cluster
- Process
- Network topology
- MPI Process
- Rank
- Message
- Message tag, source, destination
- Communicator
- MPI point-to-point communication

Torus Interconnect (Correction)



What is the advantage of torus over mesh?

- Diameter $2(\sqrt{p}/2)$
- Bisection width $2\sqrt{p}$
- Cost $2p$

Example 4 (from previous class)

```
if (myrank == 0 || myrank == 2)
    MPI_Send(arr, 20, MPI_INT, 1, 99, MPI_COMM_WORLD);
else if (myrank == 1)
{
    int count, recvarr[3][20];
    MPI_Recv(recvarr[0], 20, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG,
MPI_COMM_WORLD, &status);

    printf("Rank %d of %d received from rank %d\n", myrank, size,
status.MPI_SOURCE);

    MPI_Recv(recvarr[2], 20, MPI_INT, MPI_ANY_SOURCE, MPI_ANY_TAG,
MPI_COMM_WORLD, &status);

    printf("Rank %d of %d received from rank %d\n", myrank, size,
status.MPI_SOURCE);
}
```

Is it the
right way?

Is it the
right way?

MPI_Recv Error Handling

Fatal error in MPI_Recv: Message truncated, error stack:


MPI_Recv(200).....: MPI_Recv(buf=0x7ffdb75742b0,
count=40, MPI_INT, src=0, tag=99, MPI_COMM_WORLD,
status=0x7ffdb7574290) failed

MPIDI_CH3U_Receive_data_found(131): Message from rank 0 and tag
99 truncated; 320 bytes received but buffer size is 160

Communicator



- Object containing a group of processes
- Representative of communication domain
- Associated with a context ID (in MPICH)
- Predefined:
 - `MPI_COMM_WORLD` (when is it defined?)
 - `MPI_COMM_SELF`
- Contains a mapping from MPI process ranks to processor ids
- Memory proportional to #processes in the group

MPI_COMM_WORLD



Subcommunicator

- Logical subset
- Different contexts




MPI_COMM_SPLIT

MPI_Comm_split (MPI_Comm oldcomm, int color, int key, MPI_Comm *newcomm)

- Collective call
- Logically divides based on *color*
 - Same color processes form a group
 - Some processes may not be part of newcomm (MPI_UNDEFINED)
- Rank assignment based on *key*

Logical subsets of processes



0 → 0
2 → 1
4 → 2
...

1 → 0
3 → 1
5 → 2
...

How do you assign one color to odd processes and another color to even processes ?
 $\text{color} = \text{rank} \% 2$

Example code

```
int newrank, newsize, color = myrank%3;
MPI_Comm newcomm;

MPI_Comm_split (MPI_COMM_WORLD, color, myrank, &newcomm);

MPI_Comm_rank (newcomm, &newrank);
MPI_Comm_size (newcomm, &newsize);
printf ("%d: %d of %d\n", myrank, newrank, newsize);

MPI_Comm_free (&newcomm);
```


OUTPUT for n=9

```
0: 0 of 3
1: 0 of 3
2: 0 of 3
3: 1 of 3
4: 1 of 3
5: 1 of 3
6: 2 of 3
7: 2 of 3
8: 2 of 3
```

Process Group

- Ordered set of processes
- Ranks are contiguous
- Base group – associated with MPI_COMM_WORLD
- MPI_Group object
 - MPI_Group_rank, MPI_Group_size
- Unions and intersections of groups
- Not used in communication context

Inter and Intra Communicators



Collective Communications

- Must be called by all processes that are part of the communicator

Types

- Synchronization (MPI_Barrier)
- Global communication (MPI_Bcast, MPI_Gather, ...)
- Global reduction (MPI_Reduce, ..)

Barrier

- MPI_Barrier (comm)
- Collective call
- Caller returns only after all processes have entered the call

$n=4$

```
if (myrank != 0)
    MPI_Barrier (MPI_COMM_WORLD);
printf("%d\n", rank);
```


Output?

Broadcast

- Root process sends message to all processes
- Any process can be root process but has to be the same
- `int MPI_Bcast (buffer, count, datatype, root, comm)`
- Number of elements in buffer – count
- Tag?
- buffer – Input or output?

Q1: Can you use point-to-point communication for the same?

Q2: Buffer size of “buf” array is not known apriori at non-root processes, how should root broadcast buf?




Gather

- Gathers values from all processes to a root process
- `int MPI_Gather (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm)`
- Arguments `recv*` not relevant on non-root processes


Q: Equivalent point-to-point communications for the same?

- `MPI_Recv` at root
- `MPI_Send` at non-root

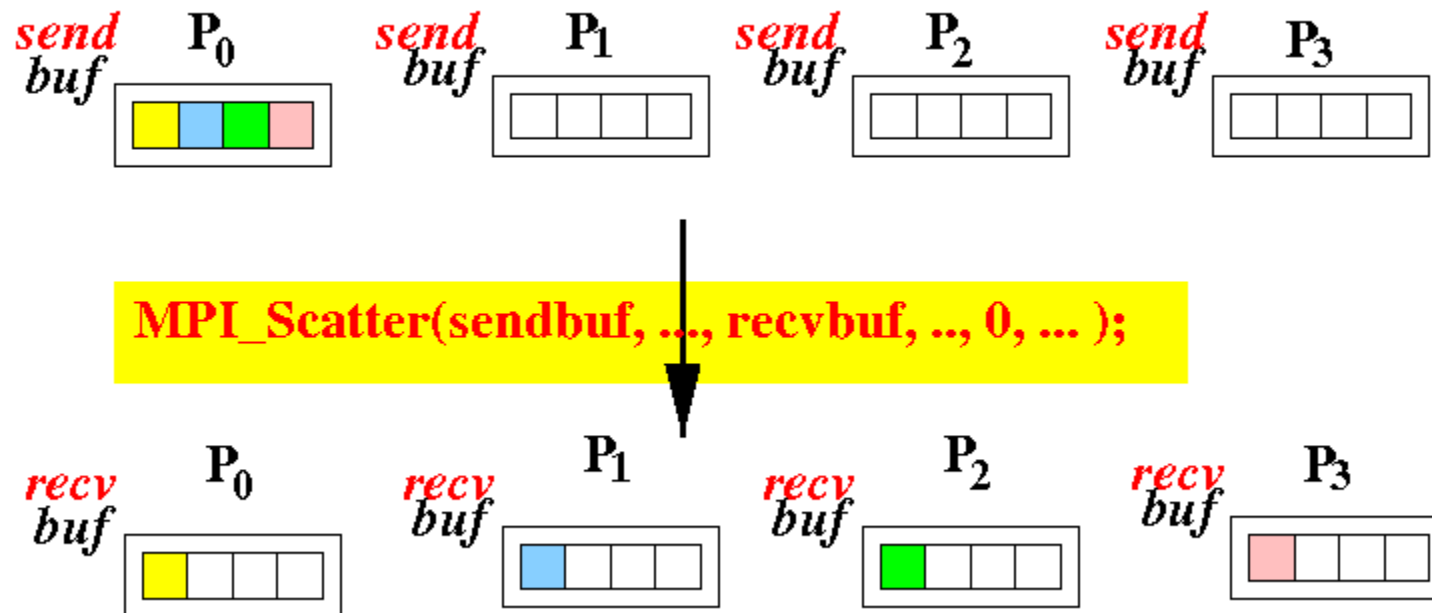


Scatter

- Scatters values to all processes from a root process
- `int MPI_Scatter (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm)`
- Arguments `send*` not relevant on non-root processes
- Output parameter – `recvbuf`




MPI_Scatter Illustration



Credit: Shun Yan Cheung

Allgather

- All processes gather values from all processes
- `int MPI_Allgather (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)`

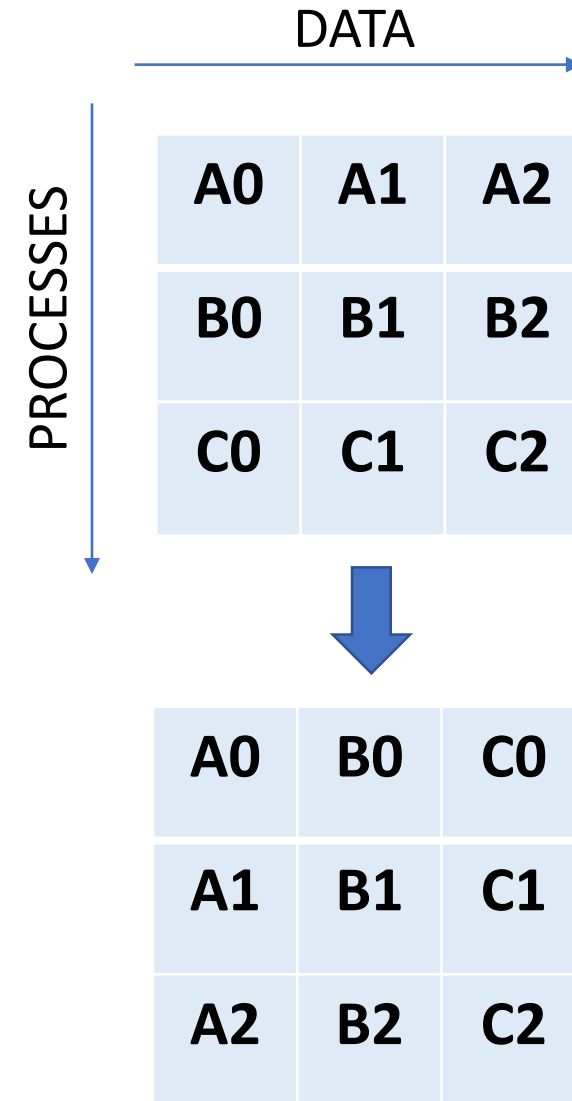


Alltoall

- Send data from all processes to all processes
- `int MPI_Alltoall (sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)`
- Output parameter – `recvbuf`

Equivalent collective?

- `MPI_Scatter` at all processes
- Cons?




Gatherv

- Root gathers values of different lengths from all processes
- `int MPI_Gatherv (sendbuf, sendcount, sendtype, recvbuf, recvcounts, displs, recvtype, root, comm)`
- `recvcounts` – Number of elements to be received from each process
- `displs` – Displacement at which to place received data


`MPI_Recv (recvbuf+displs[i], recvcounts[i], recvtype, i, i, comm, &status)` at root

`MPI_Send` at non-root




Allgatherv

- All processes gather values of different lengths from all processes
- `int MPI_Allgatherv (sendbuf, sendcount, sendtype, recvbuf, recvcunts, displs, recvtype, comm)`
- `recvcunts` – Number of elements to be received from each process
- `displs` – Displacement at which to place received data



Alltoallv

- Every process sends data of different lengths to other processes
- `int MPI_Alltoallv (sendbuf, sendcount, sdispls, sendtype, recvbuf, recvcount, rdispls, recvtype, comm)`
- Output parameter – `recvbuf`



Assignments

- Directory named AssignmentN (not assignment, Assignment- etc.)
- Reporter access at least (not guest)
- Indicate your choice of score reduction or utilizing extra day
- Start coding early
 - Parallel programs are harder to debug!
 - Your success depends on system availability

Assignment 1

1.1: Implement a modified version of the classic producer-consumer problem on distributed memory systems using N processes. You may choose P producers and C consumers, where $P+C=N$ and $P=C=N/2$. Assume that each producer produces D bytes of data (use doubles as datatype). Assume that there is a one-to-one mapping between producers and consumers.

1.2: Implement 1.1 using MPI collectives.

1.3: (Bonus question) Choose the P producer processes and $P \rightarrow C$ mapping optimally among the N processes such that performance is expected to improve.